

WHERE WE ARE

Our should-cost today: capable, but hands-on

We already cost parts properly. On the left is that job as it runs today. On the right is what it is worth, in our own numbers.

OUR SHOULD-COST TODAY · capable, but hands-on

Trusted and thorough

A rigorous bottom-up method, across every commodity.

Hands-on setup

Material, process and machine chosen by a person, part by part.

Manual data entry

Geometry and tolerances read off the 3D model and the drawings, then typed in — 12 to 69 values per part.

Updated by hand

Rates and inputs refreshed manually as economics move.

What it is worth: five numbers, four of them ours

	What it is	Our number
A	Minutes to cost one part today	_____
B	Of those, minutes reading the model and typing	_____
C	Parts we cost in a year	_____
D	Parts in a basket we never cost individually	_____
E	Share of those values the software fills itself	_____

Hours spent typing today

$$B \times C \div 60$$

Hours freed a year

$$B \times E \times C \div 60$$

Extra parts that time would cost

$$B \times E \times C \div A$$


The step forward: the same defensible should-cost, with the manual work automated

Nothing about the method changes. The judgement, the build-up and the numbers we would defend to a supplier all stay as they are. What we are proposing to remove is the typing, and then the limit on how many parts we can get through.

OPTION 1 · START NOW

Option 1: automate the data entry, inside CAPEE

CAPEE keeps doing the costing. The engineer stops typing the input and starts checking it. Reads the 3D model only, so no AI is involved.



WHAT THE 3D MODEL GIVES, WITH NOTHING TYPED

Weight, volume, overall size, surface area.
Wall thickness, holes, pockets, bosses, bends,
gear teeth, draft, machined face count.

Fixed rules turn those measurements into the values the CAPEE form asks for.

WHAT A PERSON STILL SUPPLIES

Which material it is. A 3D model cannot tell steel from aluminium, and the same shape weighs about three times more in one than the other.

**Tolerance class, finish, heat treatment, coating.
The tool asks. It never guesses and never pre-ticks.**

Why this one can start now

The 3D model is measured by ordinary engineering software, the same kind of geometry engine that sits inside a CAD package, and the values are filled in by fixed rules. There is no AI in this option, so it needs no AI approval. The file is read on a JLR server and never leaves it, and CAPEE still does the costing on the same screens.

OPTION 2 · THE CAPABILITY

Option 2: what CostVision does

Six things we cannot do at all today. All of them exist in the tool now.



Cost a whole basket, unattended

A list of parts and their 3D models goes in. Every part is costed and reported without anyone sitting there.



A board photo becomes a costed BOM

Take a picture of a circuit board. The tool works out what is on it, builds the bill of materials and the board spec, and costs it.



An RFQ pack becomes costed lines

The tool pulls the line items out of the pack, costs each one, ranks them by money and drafts the negotiation brief.



An assistant does the setting up

It picks the process route and the machine and fills in what the geometry gives, then asks one question instead of a hundred.



Supplier quotes taken apart

Put their price next to our own build-up and see which part of it the gap is in, with the questions to put to them.



Design changes priced, not guessed

It flags what makes a part expensive and puts a number on each change, ranked by the money it would save.

All six end in the same engine

The AI reads, sorts and sets up — the photo, the pack, the drawing, the process route. It never sets a price. The money is worked out by the same fixed rules and the same rate book we use today, so every number can still be explained line by line.

THE DECISION

What changes, and what we are asking for

	Today	With Option 2
How many parts get a proper cost	The ones we have hours for	All of them
Turnaround on a basket	Days of engineer time	One run, once the data is together
What we take to a supplier	Our estimate	The build-up, where their price differs, and what to ask
Where the engineer spends the time	Typing values in	Answering what only a person can, and acting on what it finds

OPTION 1 NEEDS

A Linux server, and a way to pass numbers into CAPEE
30 to 50 parts where we know the price we paid
One commodity and a small team for the trial

OPTION 2 NEEDS

Everything above, plus approval to use AI
Our own rate book loaded, with a change history
A record of every costing, and the run itself built

What we are asking for today: approve Option 1 as a trial, and start the AI approval for Option 2

One commodity, 30 to 50 parts we have already bought, inside CAPEE, using our own time. It gives us the two numbers nobody can state today: how close the tool gets to a price we actually paid, and how much of the input it fills on its own.

Two things to be straight about

Everything on the last slide works today one part at a time. Making it run over a whole basket unattended is real work, not a switch. And we have never checked the tool against a price JLR has actually paid — which is exactly what the Option 1 trial settles.