

WHERE WE ARE

Our should-cost today: capable, but hands-on

We already cost parts properly. On the left is that job as it runs today. On the right is what it is worth, in our own numbers.

OUR SHOULD-COST TODAY · capable, but hands-on

Trusted and thorough

A rigorous bottom-up method, across every commodity.

Hands-on setup

Material, process and machine chosen by a person, part by part.

Manual data entry

Geometry and tolerances read off the 3D model and the drawings, then typed in — 12 to 69 values per part.

Updated by hand

Rates and inputs refreshed manually as economics move.

What it is worth: five numbers, four of them ours

	What it is	Our number
A	Minutes to cost one part today	_____
B	Of those, minutes reading the model and typing	_____
C	Parts we cost in a year	_____
D	Parts in a basket we never cost individually	_____
E	Share of those values the software fills itself	_____

Hours spent typing today

$$B \times C \div 60$$

Hours freed a year

$$B \times E \times C \div 60$$

Extra parts that time would cost

$$B \times E \times C \div A$$


The step forward: the same defensible should-cost, with the manual work automated

Nothing about the method changes. The judgement, the build-up and the numbers we would defend to a supplier all stay as they are. What we are proposing to remove is the typing, and then the limit on how many parts we can get through.

OPTION 1 · START NOW

Option 1: automate the data entry, inside CAPEE

CAPEE keeps doing the costing. The engineer stops typing the input and starts checking it. Reads the 3D model only, so no AI is involved.



WHAT THE 3D MODEL GIVES, WITH NOTHING TYPED

Weight, volume, overall size, surface area.
Wall thickness, holes, pockets, bosses, bends,
gear teeth, draft, machined face count.

Fixed rules turn those measurements into the values the CAPEE form asks for.

WHAT A PERSON STILL SUPPLIES

Which material it is. A 3D model cannot tell steel from aluminium, and the same shape weighs about three times more in one than the other.

**Tolerance class, finish, heat treatment, coating.
The tool asks. It never guesses and never pre-ticks.**

Why this one can start now

The 3D model is measured by ordinary engineering software, the same kind of geometry engine that sits inside a CAD package, and the values are filled in by fixed rules. There is no AI in this option, so it needs no AI approval. The file is read on a JLR server and never leaves it, and CAPEE still does the costing on the same screens.

OPTION 2 · THE CAPABILITY

Option 2: should-cost that runs by itself

This is not Option 1 at a larger scale. It is three things we cannot do at all today.



Cost the whole basket, not a sample

A list of parts and their 3D models goes in. Every part is costed from its own geometry, and a report comes out for each part and for the basket. Today we cost what we have hours for and estimate the rest from a sample.



An assistant does the setting up

An AI assistant picks the process route and the machine, fills in what the geometry gives, and comes back with one question instead of a hundred — answered once and applied to every part it affects. It never sets a price: the money is worked out by the same fixed rules as today.



Answers we cannot get today

Put a supplier quote next to our own build-up and see which part of it the gap is in, with the questions to ask them. Savings ideas priced rather than just listed. Where a part should be made, ten countries side by side. What happens if the metal price or the volume moves. And the tool narrows its own range as we feed it prices we have actually paid.

FIRST WE NEED

Approval to use AI

Our own rate book

A history of rate changes

A record of every costing

The run itself built

THE DECISION

What it is worth, and what we are asking for

OPTION 1 NEEDS

A Linux server, and a way to pass numbers into CAPEE
30 to 50 parts where we know the price we paid
One part type and a small team for the trial

OPTION 2 NEEDS

Everything above, plus approval to use AI
Our own rate book loaded, with a change history
A record of every costing, and the run itself built

What it is worth: five numbers, four of them yours

	What it is	Our number	Where it comes from
A	Minutes to cost one part today	_____	Time a handful of parts. One commodity is enough to start
B	Of those, the minutes spent reading the model and typing values in	_____	The same exercise, timed separately
C	Parts we cost in a year	_____	Our own records
D	Parts in a typical basket we never cost individually	_____	Programme data. This is the ground Option 2 covers
E	Share of those values the software fills on its own	_____	The trial measures this, per commodity

Hours spent typing today

$$B \times C \div 60$$

Hours freed a year

$$B \times E \times C \div 60$$

Extra parts that time would cost

$$B \times E \times C \div A$$

What we are asking for today: approve Option 1 as a trial, and start the AI approval for Option 2

One commodity, 30 to 50 parts we have already bought, inside CAPEE, using our own time. It gives us the two numbers nobody can state today: how close the tool gets to a price we actually paid, and how much of the input it fills on its own.